

Normalising flows with applications

ML Masterclass

David Huk and Rigers Behluli

University of Warwick

May 22, 2023

Summary

- 1 Introduction
- 2 Basics
- 3 Applications
- 4 Methods
- 5 Computer simulations
- 6 References

Summary

- 1 Introduction
- 2 Basics
- 3 Applications
- 4 Methods
- 5 Computer simulations
- 6 References

Task: model probability distribution from data

Existing methods:

- Direct analytical approaches
- Variational approaches
- Expectation maximization
- VAEs and GANs

VAEs and GANs do not allow for likelihood evaluation and may present some training issues (e.g. mode collapse, posterior collapse, vanishing gradients and training instability).

→ Normalizing Flows ([Kobyzev et al., 2020], review paper)

Introduction

Task: model probability distribution from data

Existing methods:

- Direct analytical approaches
- Variational approaches
- Expectation maximization
- VAEs and GANs

VAEs and GANs do not allow for likelihood evaluation and may present some training issues (e.g. mode collapse, posterior collapse, vanishing gradients and training instability).

→ Normalizing Flows ([Kobyzev et al., 2020], review paper)

Introduction

Task: model probability distribution from data

Existing methods:

- Direct analytical approaches
- Variational approaches
- Expectation maximization
- VAEs and GANs

VAEs and GANs do not allow for likelihood evaluation and may present some training issues (e.g. mode collapse, posterior collapse, vanishing gradients and training instability).

→ Normalizing Flows ([Kobyzev et al., 2020], review paper)

Introduction

Task: model probability distribution from data

Existing methods:

- Direct analytical approaches
- Variational approaches
- Expectation maximization
- VAEs and GANs

VAEs and GANs do not allow for likelihood evaluation and may present some training issues (e.g. mode collapse, posterior collapse, vanishing gradients and training instability).

→ Normalizing Flows ([Kobyzev et al., 2020], review paper)

Introduction

Task: model probability distribution from data

Existing methods:

- Direct analytical approaches
- Variational approaches
- Expectation maximization
- VAEs and GANs

VAEs and GANs do not allow for likelihood evaluation and may present some training issues (e.g. mode collapse, posterior collapse, vanishing gradients and training instability).

→ Normalizing Flows ([Kobyzev et al., 2020], review paper)

Introduction

Task: model probability distribution from data

Existing methods:

- Direct analytical approaches
- Variational approaches
- Expectation maximization
- VAEs and GANs

VAEs and GANs do not allow for likelihood evaluation and may present some training issues (e.g. mode collapse, posterior collapse, vanishing gradients and training instability).

→ Normalizing Flows ([Kobyzev et al., 2020], review paper)

Task: model probability distribution from data

Existing methods:

- Direct analytical approaches
- Variational approaches
- Expectation maximization
- VAEs and GANs

VAEs and GANs do not allow for likelihood evaluation and may present some training issues (e.g. mode collapse, posterior collapse, vanishing gradients and training instability).

→ Normalizing Flows ([Kobyzev et al., 2020], review paper)

Summary

- 1 Introduction
- 2 Basics**
- 3 Applications
- 4 Methods
- 5 Computer simulations
- 6 References

Normalising flows basics

Let $\mathbf{Z} \in \mathbb{R}^D$ be a random variable with a known probability density function $p_{\mathbf{Z}} : \mathbb{R} \rightarrow \mathbb{R}$ and let \mathbf{g} be an invertible function and $\mathbf{Y} = \mathbf{g}(\mathbf{Z})$. Using the change of variables formula, one can compute:

$$p_{\mathbf{Y}}(\mathbf{y}) = p_{\mathbf{Z}}(\mathbf{f}(\mathbf{y})) |\det D\mathbf{g}(\mathbf{f}(\mathbf{y}))|^{-1},$$

where \mathbf{f} is the inverse of \mathbf{g} and $\det D\mathbf{g}(\mathbf{z}) = \frac{\partial \mathbf{g}}{\partial \mathbf{z}}$ is the Jacobian of g .

The function \mathbf{g} transforms the the base density $p_{\mathbf{Z}}$ into a more complex density, defining also it's generative direction.

The inverse of \mathbf{g} instead moves in the opposite direction, from a complicated density to a know one.

Normalising flows basics

Let $\mathbf{Z} \in \mathbb{R}^D$ be a random variable with a known probability density function $p_{\mathbf{Z}} : \mathbb{R} \rightarrow \mathbb{R}$ and let \mathbf{g} be an invertible function and $\mathbf{Y} = \mathbf{g}(\mathbf{Z})$. Using the change of variables formula, one can compute:

$$p_{\mathbf{Y}}(\mathbf{y}) = p_{\mathbf{Z}}(\mathbf{f}(\mathbf{y})) |\det D\mathbf{g}(\mathbf{f}(\mathbf{y}))|^{-1},$$

where \mathbf{f} is the inverse of \mathbf{g} and $\det D\mathbf{g}(\mathbf{z}) = \frac{\partial \mathbf{g}}{\partial \mathbf{z}}$ is the Jacobian of g .

The function \mathbf{g} transforms the the base density $p_{\mathbf{Z}}$ into a more complex density, defining also it's generative direction.

The inverse of \mathbf{g} instead moves in the opposite direction, from a complicated density to a know one.

Normalising flows basics

Let $\mathbf{Z} \in \mathbb{R}^D$ be a random variable with a known probability density function $p_{\mathbf{Z}} : \mathbb{R} \rightarrow \mathbb{R}$ and let \mathbf{g} be an invertible function and $\mathbf{Y} = \mathbf{g}(\mathbf{Z})$. Using the change of variables formula, one can compute:

$$p_{\mathbf{Y}}(\mathbf{y}) = p_{\mathbf{Z}}(\mathbf{f}(\mathbf{y})) |\det D\mathbf{g}(\mathbf{f}(\mathbf{y}))|^{-1},$$

where \mathbf{f} is the inverse of \mathbf{g} and $\det D\mathbf{g}(\mathbf{z}) = \frac{\partial \mathbf{g}}{\partial \mathbf{z}}$ is the Jacobian of g .

The function \mathbf{g} transforms the the base density $p_{\mathbf{Z}}$ into a more complex density, defining also it's generative direction.

The inverse of \mathbf{g} instead moves in the opposite direction, from a complicated density to a know one.

Normalising flows basics

Let $\mathbf{Z} \in \mathbb{R}^D$ be a random variable with a known probability density function $p_{\mathbf{Z}} : \mathbb{R} \rightarrow \mathbb{R}$ and let \mathbf{g} be an invertible function and $\mathbf{Y} = \mathbf{g}(\mathbf{Z})$. Using the change of variables formula, one can compute:

$$p_{\mathbf{Y}}(\mathbf{y}) = p_{\mathbf{Z}}(\mathbf{f}(\mathbf{y})) |\det D\mathbf{g}(\mathbf{f}(\mathbf{y}))|^{-1},$$

where \mathbf{f} is the inverse of \mathbf{g} and $\det D\mathbf{g}(\mathbf{z}) = \frac{\partial \mathbf{g}}{\partial \mathbf{z}}$ is the Jacobian of g .

The function \mathbf{g} transforms the the base density $p_{\mathbf{Z}}$ into a more complex density, defining also it's generative direction.

The inverse of \mathbf{g} instead moves in the opposite direction, from a complicated density to a know one.

Summary

- 1 Introduction
- 2 Basics
- 3 Applications**
- 4 Methods
- 5 Computer simulations
- 6 References

Application - Density estimation

- Data $y_1, \dots, y_n \sim p_Y$
- Base noise $z \sim p_Z(\cdot; \phi)$
- Flow model $y = \mathbf{g}_\theta(z)$ with inverse \mathbf{f}_θ

$$\begin{aligned}\log p(y_1, \dots, y_n \mid \theta, \phi) &= \sum_{i=1}^n \log p_Y(y_i \mid \theta, \phi) \\ &= \sum_{i=1}^n \underbrace{\log p_Z(\mathbf{f}(y_i \mid \theta) \mid \phi)}_{\text{ll under base measure}} + \underbrace{\log |\det D\mathbf{f}(y_i \mid \theta)|}_{\text{volume correction}}\end{aligned}$$

During training, estimate θ, ϕ to maximise the above ll.

Computational cost depends on computing \mathbf{f} and $\det D\mathbf{f}$ - the *normalising direction*.

Application - sampling

To sample from p_Y , it is enough to generate noise $z \sim p_Z$ and apply the flow \mathbf{g} .

Performance determined by the cost of applying \mathbf{g} - the *generative direction*.




Figure: Generative direction (sampling) and Normalising direction (likelihood).

In general, want **invertible** and **expressive** \mathbf{g} with **efficient evaluation** of \mathbf{g}, \mathbf{f} and $\det D\mathbf{f}$. Different flows have different strengths.

Summary

- 1 Introduction
- 2 Basics
- 3 Applications
- 4 Methods**
- 5 Computer simulations
- 6 References

Elementwise Flows

Normalizing flows should satisfy several conditions to be practical:

- be invertible,
- be sufficiently expressive to model the distribution of interest,
- be computationally efficient.

An example could be to use any bijective scalar function.

Let $h : \mathbb{R} \rightarrow \mathbb{R}$ be a scalar valued bijection. If $\mathbf{x} = (x_1, x_2, \dots, x_D)^T$, then

$$\mathbf{g}(\mathbf{x}) = (h(x_1), h(x_2), \dots, h(x_D))^T$$

is also a bijection whose inverse simply requires computing h^{-1} and whose Jacobian is the product of the absolute values of the derivatives of h .

Very simple approach, but limiting.

Elementwise Flows

Normalizing flows should satisfy several conditions to be practical:

- be invertible,
- be sufficiently expressive to model the distribution of interest,
- be computationally efficient.

An example could be to use any bijective scalar function.

Let $h : \mathbb{R} \rightarrow \mathbb{R}$ be a scalar valued bijection. If $\mathbf{x} = (x_1, x_2, \dots, x_D)^T$, then

$$\mathbf{g}(\mathbf{x}) = (h(x_1), h(x_2), \dots, h(x_D))^T$$

is also a bijection whose inverse simply requires computing h^{-1} and whose Jacobian is the product of the absolute values of the derivatives of h .

Very simple approach, but limiting.

Elementwise Flows

Normalizing flows should satisfy several conditions to be practical:

- be invertible,
- be sufficiently expressive to model the distribution of interest,
- be computationally efficient.

An example could be to use any bijective scalar function.

Let $h : \mathbb{R} \rightarrow \mathbb{R}$ be a scalar valued bijection. If $\mathbf{x} = (x_1, x_2, \dots, x_D)^T$, then

$$\mathbf{g}(\mathbf{x}) = (h(x_1), h(x_2), \dots, h(x_D))^T$$

is also a bijection whose inverse simply requires computing h^{-1} and whose Jacobian is the product of the absolute values of the derivatives of h .

Very simple approach, but limiting.

Elementwise Flows

Normalizing flows should satisfy several conditions to be practical:

- be invertible,
- be sufficiently expressive to model the distribution of interest,
- be computationally efficient.

An example could be to use any bijective scalar function.

Let $h : \mathbb{R} \rightarrow \mathbb{R}$ be a scalar valued bijection. If $\mathbf{x} = (x_1, x_2, \dots, x_D)^T$, then

$$\mathbf{g}(\mathbf{x}) = (h(x_1), h(x_2), \dots, h(x_D))^T$$

is also a bijection whose inverse simply requires computing h^{-1} and whose Jacobian is the product of the absolute values of the derivatives of h .

Very simple approach, but limiting.

Normalizing flows should satisfy several conditions to be practical:

- be invertible,
- be sufficiently expressive to model the distribution of interest,
- be computationally efficient.

An example could be to use any bijective scalar function.

Let $h : \mathbb{R} \rightarrow \mathbb{R}$ be a scalar valued bijection. If $\mathbf{x} = (x_1, x_2, \dots, x_D)^T$, then

$$\mathbf{g}(\mathbf{x}) = (h(x_1), h(x_2), \dots, h(x_D))^T$$

is also a bijection whose inverse simply requires computing h^{-1} and whose Jacobian is the product of the absolute values of the derivatives of h .

Very simple approach, but limiting.

Normalizing flows should satisfy several conditions to be practical:

- be invertible,
- be sufficiently expressive to model the distribution of interest,
- be computationally efficient.

An example could be to use any bijective scalar function.

Let $h : \mathbb{R} \rightarrow \mathbb{R}$ be a scalar valued bijection. If $\mathbf{x} = (x_1, x_2, \dots, x_D)^T$, then

$$\mathbf{g}(\mathbf{x}) = (h(x_1), h(x_2), \dots, h(x_D))^T$$

is also a bijection whose inverse simply requires computing h^{-1} and whose Jacobian is the product of the absolute values of the derivatives of h .

Very simple approach, but limiting.

Linear Flows

Linear mappings can express correlation between dimensions.

$$\mathbf{g}(\mathbf{x}) = \mathbf{Ax} + \mathbf{b}$$

where A and b are parameters. If A is invertible, then the function is invertible too. Determinant of the Jacobian can be computed in $\mathcal{O}(D^3)$.

A can be:

- diagonal: close to elementwise flows,
- upper triangular: captures correlation between dimensions and can be computed in $\mathcal{O}(D^2)$,
- permutation and orthogonal matrices can be used for faster computations,
- factorization: can be computed in $\mathcal{O}(D)$

$$\mathbf{g}(\mathbf{x}) = \mathbf{PLUx} + \mathbf{b}$$

Linear Flows

Linear mappings can express correlation between dimensions.

$$\mathbf{g}(\mathbf{x}) = \mathbf{Ax} + \mathbf{b}$$

where A and b are parameters. If A is invertible, then the function is invertible too. Determinant of the Jacobian can be computed in $\mathcal{O}(D^3)$.

A can be:

- diagonal: close to elementwise flows,
- upper triangular: captures correlation between dimensions and can be computed in $\mathcal{O}(D^2)$,
- permutation and orthogonal matrices can be used for faster computations,
- factorization: can be computed in $\mathcal{O}(D)$

$$\mathbf{g}(\mathbf{x}) = \mathbf{PLUx} + \mathbf{b}$$

Linear Flows

Linear mappings can express correlation between dimensions.

$$\mathbf{g}(\mathbf{x}) = \mathbf{Ax} + \mathbf{b}$$

where A and b are parameters. If A is invertible, then the function is invertible too. Determinant of the Jacobian can be computed in $\mathcal{O}(D^3)$.

A can be:

- diagonal: close to elementwise flows,
- upper triangular: captures correlation between dimensions and can be computed in $\mathcal{O}(D^2)$,
- permutation and orthogonal matrices can be used for faster computations,
- factorization: can be computed in $\mathcal{O}(D)$

$$\mathbf{g}(\mathbf{x}) = \mathbf{PLUx} + \mathbf{b}$$

Linear Flows

Linear mappings can express correlation between dimensions.

$$\mathbf{g}(\mathbf{x}) = \mathbf{Ax} + \mathbf{b}$$

where A and b are parameters. If A is invertible, then the function is invertible too. Determinant of the Jacobian can be computed in $\mathcal{O}(D^3)$.

A can be:

- diagonal: close to elementwise flows,
- upper triangular: captures correlation between dimensions and can be computed in $\mathcal{O}(D^2)$,
- permutation and orthogonal matrices can be used for faster computations,
- factorization: can be computed in $\mathcal{O}(D)$

$$\mathbf{g}(\mathbf{x}) = \mathbf{PLUx} + \mathbf{b}$$

Linear Flows

Linear mappings can express correlation between dimensions.

$$\mathbf{g}(\mathbf{x}) = \mathbf{Ax} + \mathbf{b}$$

where A and b are parameters. If A is invertible, then the function is invertible too. Determinant of the Jacobian can be computed in $\mathcal{O}(D^3)$.

A can be:

- diagonal: close to elementwise flows,
- upper triangular: captures correlation between dimensions and can be computed in $\mathcal{O}(D^2)$,
- permutation and orthogonal matrices can be used for faster computations,
- factorization: can be computed in $\mathcal{O}(D)$

$$\mathbf{g}(\mathbf{x}) = \mathbf{PLUx} + \mathbf{b}$$

Coupling Flows (1)

Introduced by [Dinh et al., 2014], the idea is to only transform subsets of data at a time.

- Split input $\mathbf{x} \in \mathbb{R}^D$ into $\mathbf{x}^A \in \mathbb{R}^d$ and $\mathbf{x}^B \in \mathbb{R}^{D-d}$.
- *Coupling function* is a bijection $\mathbf{h}(\cdot, \theta) : \mathbb{R}^d \rightarrow \mathbb{R}^d$.
- *Conditioner* $\Theta(\cdot)$ used to define θ given inputs.

Define the *coupling flow* $\mathbf{g} : \mathbb{R}^D \rightarrow \mathbb{R}^D$ as:

$$\begin{aligned}\mathbf{y}^A &= \mathbf{h}(\mathbf{x}^A, \theta = \Theta(\mathbf{x}^B)) \\ \mathbf{y}^B &= \mathbf{x}^B\end{aligned}$$

Coupling Flows (1)

Introduced by [Dinh et al., 2014], the idea is to only transform subsets of data at a time.

- Split input $\mathbf{x} \in \mathbb{R}^D$ into $\mathbf{x}^A \in \mathbb{R}^d$ and $\mathbf{x}^B \in \mathbb{R}^{D-d}$.
- *Coupling function* is a bijection $\mathbf{h}(\cdot, \theta) : \mathbb{R}^d \rightarrow \mathbb{R}^d$.
- *Conditioner* $\Theta(\cdot)$ used to define θ given inputs.

Define the *coupling flow* $\mathbf{g} : \mathbb{R}^D \rightarrow \mathbb{R}^D$ as:

$$\begin{aligned}\mathbf{y}^A &= \mathbf{h}(\mathbf{x}^A, \theta = \Theta(\mathbf{x}^B)) \\ \mathbf{y}^B &= \mathbf{x}^B\end{aligned}$$

Coupling Flows (1)

Introduced by [Dinh et al., 2014], the idea is to only transform subsets of data at a time.

- Split input $\mathbf{x} \in \mathbb{R}^D$ into $\mathbf{x}^A \in \mathbb{R}^d$ and $\mathbf{x}^B \in \mathbb{R}^{D-d}$.
- *Coupling function* is a bijection $\mathbf{h}(\cdot, \theta) : \mathbb{R}^d \rightarrow \mathbb{R}^d$.
- *Conditioner* $\Theta(\cdot)$ used to define θ given inputs.

Define the *coupling flow* $\mathbf{g} : \mathbb{R}^D \rightarrow \mathbb{R}^D$ as:

$$\begin{aligned}\mathbf{y}^A &= \mathbf{h}(\mathbf{x}^A, \theta = \Theta(\mathbf{x}^B)) \\ \mathbf{y}^B &= \mathbf{x}^B\end{aligned}$$

Coupling Flows (1)

Introduced by [Dinh et al., 2014], the idea is to only transform subsets of data at a time.

- Split input $\mathbf{x} \in \mathbb{R}^D$ into $\mathbf{x}^A \in \mathbb{R}^d$ and $\mathbf{x}^B \in \mathbb{R}^{D-d}$.
- *Coupling function* is a bijection $\mathbf{h}(\cdot, \theta) : \mathbb{R}^d \rightarrow \mathbb{R}^d$.
- *Conditioner* $\Theta(\cdot)$ used to define θ given inputs.

Define the *coupling flow* $\mathbf{g} : \mathbb{R}^D \rightarrow \mathbb{R}^D$ as:

$$\begin{aligned} \mathbf{y}^A &= \mathbf{h}(\mathbf{x}^A, \theta = \Theta(\mathbf{x}^B)) \\ \mathbf{y}^B &= \mathbf{x}^B \end{aligned}$$

Coupling Flows (1)

Introduced by [Dinh et al., 2014], the idea is to only transform subsets of data at a time.

- Split input $\mathbf{x} \in \mathbb{R}^D$ into $\mathbf{x}^A \in \mathbb{R}^d$ and $\mathbf{x}^B \in \mathbb{R}^{D-d}$.
- *Coupling function* is a bijection $\mathbf{h}(\cdot, \theta) : \mathbb{R}^d \rightarrow \mathbb{R}^d$.
- *Conditioner* $\Theta(\cdot)$ used to define θ given inputs.

Define the *coupling flow* $\mathbf{g} : \mathbb{R}^D \rightarrow \mathbb{R}^D$ as:

$$\begin{aligned}\mathbf{y}^A &= \mathbf{h}(\mathbf{x}^A, \theta = \Theta(\mathbf{x}^B)) \\ \mathbf{y}^B &= \mathbf{x}^B\end{aligned}$$

Coupling Flows (2)

The Jacobian of \mathbf{g} is a block triangular matrix where diagonal blocks are $D\mathbf{h}$ and the identity. So the determinant of the Jacobian is $\det D\mathbf{h}$.




Figure: A single coupling flow illustration from [Kobyzev et al., 2020].

The power of coupling flows is in the ability to have very complex models for $\Theta(\mathbf{x}^B)$ (which do not need to be invertible), such as neural networks.

Coupling Flows (2)

The Jacobian of \mathbf{g} is a block triangular matrix where diagonal blocks are $D\mathbf{h}$ and the identity. So the determinant of the Jacobian is $\det D\mathbf{h}$.




Figure: A single coupling flow illustration from [Kobyzev et al., 2020].

The power of coupling flows is in the ability to have very complex models for $\Theta(\mathbf{x}^B)$ (which do not need to be invertible), such as neural networks.

Autoregressive Flows (1)

Similar to coupling flows, but with each entry of outputs $\mathbf{y} = \mathbf{g}(\mathbf{x})$ conditioned on the previous entries of the input.

- *Coupling function* is a bijection $h(., \theta) : \mathbb{R} \rightarrow \mathbb{R}$.
- *Conditioner* $\Theta(.)$ used to define θ given inputs.

Then the AR flow is given by:

$$y_t = h(x_t; \Theta_t(\mathbf{x}_{1:t-1})), \quad t \in \{1, \dots, D\}.$$

The Jacobian of \mathbf{g} is triangular with inputs y_t depending only on $\mathbf{x}_{1:t}$, so the determinant is a product of its diagonal entries.

Furthermore, \mathbf{g} can be parallelised easily due to a dependence on available subsets of \mathbf{x} only.

Autoregressive Flows (1)

Similar to coupling flows, but with each entry of outputs $\mathbf{y} = \mathbf{g}(\mathbf{x})$ conditioned on the previous entries of the input.

- *Coupling function* is a bijection $h(., \theta) : \mathbb{R} \rightarrow \mathbb{R}$.
- *Conditioner* $\Theta(.)$ used to define θ given inputs.

Then the AR flow is given by:

$$y_t = h(x_t; \Theta_t(\mathbf{x}_{1:t-1})), \quad t \in \{1, \dots, D\}.$$

The Jacobian of \mathbf{g} is triangular with inputs y_t depending only on $\mathbf{x}_{1:t}$, so the determinant is a product of its diagonal entries.

Furthermore, \mathbf{g} can be parallelised easily due to a dependence on available subsets of \mathbf{x} only.

Autoregressive Flows (1)

Similar to coupling flows, but with each entry of outputs $\mathbf{y} = \mathbf{g}(\mathbf{x})$ conditioned on the previous entries of the input.

- *Coupling function* is a bijection $h(., \theta) : \mathbb{R} \rightarrow \mathbb{R}$.
- *Conditioner* $\Theta(.)$ used to define θ given inputs.

Then the AR flow is given by:

$$y_t = h(x_t; \Theta_t(\mathbf{x}_{1:t-1})), \quad t \in \{1, \dots, D\}.$$

The Jacobian of \mathbf{g} is triangular with inputs y_t depending only on $\mathbf{x}_{1:t}$, so the determinant is a product of its diagonal entries.

Furthermore, \mathbf{g} can be parallelised easily due to a dependence on available subsets of \mathbf{x} only.

Autoregressive Flows (1)

Similar to coupling flows, but with each entry of outputs $\mathbf{y} = \mathbf{g}(\mathbf{x})$ conditioned on the previous entries of the input.

- *Coupling function* is a bijection $h(., \theta) : \mathbb{R} \rightarrow \mathbb{R}$.
- *Conditioner* $\Theta(.)$ used to define θ given inputs.

Then the AR flow is given by:

$$y_t = h(x_t; \Theta_t(\mathbf{x}_{1:t-1})), \quad t \in \{1, \dots, D\}.$$

The Jacobian of \mathbf{g} is triangular with inputs y_t depending only on $x_{1:t}$, so the determinant is a product of its diagonal entries.

Furthermore, \mathbf{g} can be parallelised easily due to a dependence on available subsets of \mathbf{x} only.

Autoregressive Flows (1)

Similar to coupling flows, but with each entry of outputs $\mathbf{y} = \mathbf{g}(\mathbf{x})$ conditioned on the previous entries of the input.

- *Coupling function* is a bijection $h(., \theta) : \mathbb{R} \rightarrow \mathbb{R}$.
- *Conditioner* $\Theta(.)$ used to define θ given inputs.

Then the AR flow is given by:

$$y_t = h(x_t; \Theta_t(\mathbf{x}_{1:t-1})), \quad t \in \{1, \dots, D\}.$$

The Jacobian of \mathbf{g} is triangular with inputs y_t depending only on $x_{1:t}$, so the determinant is a product of its diagonal entries.

Furthermore, \mathbf{g} can be parallelised easily due to a dependence on available subsets of \mathbf{x} only.

Autoregressive Flows (1)

Similar to coupling flows, but with each entry of outputs $\mathbf{y} = \mathbf{g}(\mathbf{x})$ conditioned on the previous entries of the input.

- *Coupling function* is a bijection $h(., \theta) : \mathbb{R} \rightarrow \mathbb{R}$.
- *Conditioner* $\Theta(.)$ used to define θ given inputs.

Then the AR flow is given by:

$$y_t = h(x_t; \Theta_t(\mathbf{x}_{1:t-1})), \quad t \in \{1, \dots, D\}.$$

The Jacobian of \mathbf{g} is triangular with inputs y_t depending only on $x_{1:t}$, so the determinant is a product of its diagonal entries.

Furthermore, \mathbf{g} can be parallelised easily due to a dependence on available subsets of \mathbf{x} only.

Autoregressive Flows (2)

However, the inverse flow \mathbf{f} is more complicated and is found recursively, making it hard to implement efficiently.

An alternative is proposed by [Kingma et al., 2016] under the name of *inverse autoregressive flows* (IAF), which outputs terms from \mathbf{y} conditioned on previous entries of \mathbf{y} :

$$y_t = h(x_t; \Theta_t(\mathbf{y}_{1:t-1}))$$

While most flows are modeled in the normalising direction to ensure efficient evaluations of \mathbf{f} , the IAF does the opposite and models the generative direction.

- Classical AR flows have efficient density estimation.
- IAFs have efficient sampling.

Autoregressive Flows (2)

However, the inverse flow \mathbf{f} is more complicated and is found recursively, making it hard to implement efficiently.

An alternative is proposed by [Kingma et al., 2016] under the name of *inverse autoregressive flows* (IAF), which outputs terms from \mathbf{y} conditioned on previous entries of \mathbf{y} :

$$y_t = h(x_t; \Theta_t(\mathbf{y}_{1:t-1}))$$

While most flows are modeled in the normalising direction to ensure efficient evaluations of \mathbf{f} , the IAF does the opposite and models the generative direction.

- Classical AR flows have efficient density estimation.
- IAFs have efficient sampling.

Autoregressive Flows (2)

However, the inverse flow \mathbf{f} is more complicated and is found recursively, making it hard to implement efficiently.

An alternative is proposed by [Kingma et al., 2016] under the name of *inverse autoregressive flows* (IAF), which outputs terms from \mathbf{y} conditioned on previous entries of \mathbf{y} :

$$y_t = h(x_t; \Theta_t(\mathbf{y}_{1:t-1}))$$

While most flows are modeled in the normalising direction to ensure efficient evaluations of \mathbf{f} , the IAF does the opposite and models the generative direction.

- Classical AR flows have efficient density estimation.
- IAFs have efficient sampling.

Autoregressive Flows (2)

However, the inverse flow \mathbf{f} is more complicated and is found recursively, making it hard to implement efficiently.

An alternative is proposed by [Kingma et al., 2016] under the name of *inverse autoregressive flows* (IAF), which outputs terms from \mathbf{y} conditioned on previous entries of \mathbf{y} :

$$y_t = h(x_t; \Theta_t(\mathbf{y}_{1:t-1}))$$

While most flows are modeled in the normalising direction to ensure efficient evaluations of \mathbf{f} , the IAF does the opposite and models the generative direction.

- Classical AR flows have efficient density estimation.
- IAFs have efficient sampling.

Autoregressive Flows (3)





Figure: Left: classical Autoregressive Flows. Right: Inverse Autoregressive FLOws.

Summary

- 1 Introduction
- 2 Basics
- 3 Applications
- 4 Methods
- 5 Computer simulations**
- 6 References

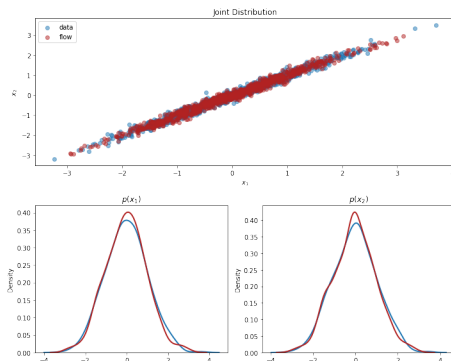
Learning a multivariate Gaussian

- Base noise is $\mathcal{N}_2(\mathbf{0}, \mathbb{I}_2)$.
- Target density is $\mathcal{N}_2(\mathbf{0}, \begin{bmatrix} 1 & 0.99 \\ 0.99 & 1 \end{bmatrix})$.
- Use a coupling flow with affine h as
$$\begin{cases} y_1 = h(x_1; \theta) = \theta_1 x_1 + \theta_2 \\ y_2 = x_2 \end{cases}$$



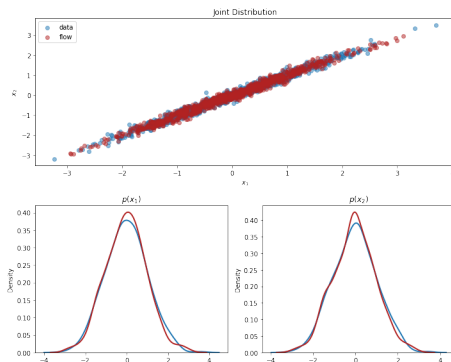
Learning a multivariate Gaussian

- Base noise is $\mathcal{N}_2(\mathbf{0}, \mathbb{I}_2)$.
- Target density is $\mathcal{N}_2(\mathbf{0}, \begin{bmatrix} 1 & 0.99 \\ 0.99 & 1 \end{bmatrix})$.
- Use a coupling flow with affine h as
$$\begin{cases} y_1 = h(x_1; \theta) = \theta_1 x_1 + \theta_2 \\ y_2 = x_2 \end{cases}$$



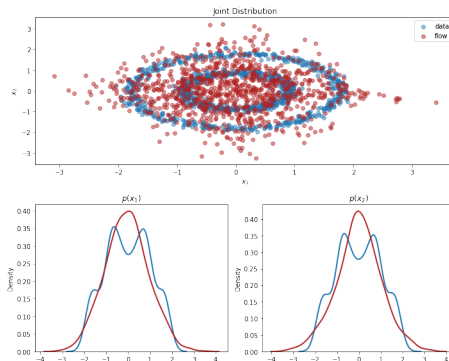
Learning a multivariate Gaussian

- Base noise is $\mathcal{N}_2(\mathbf{0}, \mathbb{I}_2)$.
- Target density is $\mathcal{N}_2(\mathbf{0}, \begin{bmatrix} 1 & 0.99 \\ 0.99 & 1 \end{bmatrix})$.
- Use a coupling flow with affine h as
$$\begin{cases} y_1 = h(x_1; \theta) = \theta_1 x_1 + \theta_2 \\ y_2 = x_2 \end{cases}$$




Learning a non-Gaussian

- Base noise is $\mathcal{N}_2(\mathbf{0}, \mathbb{I}_2)$.
- Target density is uniform on two circles.
- Use a coupling flow with affine h as
$$\begin{cases} y_1 = h(x_1; \theta) = \theta_1 x_1 + \theta_2 \\ y_2 = x_2 \end{cases}$$



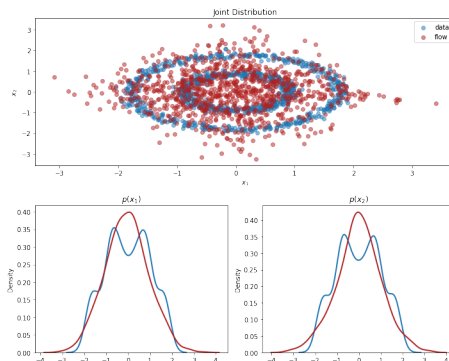
Learning a non-Gaussian

- Base noise is $\mathcal{N}_2(\mathbf{0}, \mathbb{I}_2)$.
- Target density is uniform on two circles.
- Use a coupling flow with affine h as
$$\begin{cases} y_1 = h(x_1; \theta) = \theta_1 x_1 + \theta_2 \\ y_2 = x_2 \end{cases}$$




Learning a non-Gaussian

- Base noise is $\mathcal{N}_2(\mathbf{0}, \mathbb{I}_2)$.
- Target density is uniform on two circles.
- Use a coupling flow with affine h as
$$\begin{cases} y_1 = h(x_1; \theta) = \theta_1 x_1 + \theta_2 \\ y_2 = x_2 \end{cases}$$




Learning a non-Gaussian

- Base noise is $\mathcal{N}_2(\mathbf{0}, \mathbb{I}_2)$.
- Target density is uniform on two circles.
- Use a coupling flow with spline h as
$$\begin{cases} y_1 = h(x_1; \Theta(x_2)) = s_\theta(x_1; x_2) \\ y_2 = x_2 \end{cases}$$




Learning a non-Gaussian

- Base noise is $\mathcal{N}_2(\mathbf{0}, \mathbb{I}_2)$.
- Target density is uniform on two circles.
- Use a coupling flow with spline h as
$$\begin{cases} y_1 = h(x_1; \Theta(x_2)) = s_\theta(x_1; x_2) \\ y_2 = x_2 \end{cases}$$



Learning a non-Gaussian

- Base noise is $\mathcal{N}_2(\mathbf{0}, \mathbb{I}_2)$.
- Target density is uniform on two circles.
- Use a coupling flow with spline h as
$$\begin{cases} y_1 = h(x_1; \Theta(x_2)) = s_\theta(x_1; x_2) \\ y_2 = x_2 \end{cases}$$



Summary

- 1 Introduction
- 2 Basics
- 3 Applications
- 4 Methods
- 5 Computer simulations
- 6 References**

References



Dinh, L., Krueger, D., and Bengio, Y. (2014).
Nice: Non-linear independent components estimation.
arXiv preprint arXiv:1410.8516.



Kingma, D. P., Salimans, T., Jozefowicz, R., Chen, X., Sutskever, I.,
and Welling, M. (2016).
Improved variational inference with inverse autoregressive flow.
Advances in neural information processing systems, 29.



Kobyzev, I., Prince, S. J., and Brubaker, M. A. (2020).
Normalizing flows: An introduction and review of current methods.
IEEE transactions on pattern analysis and machine intelligence,
43(11):3964–3979.

Thank you!

David Huk and Rigers Behluli